The OPAQUE Asymmetric PAKE Protocol
draft-krawczyk-cfrg-opaque-03

Table of Contents

Abstract

This draft describes the OPAQUE protocol, a secure asymmetric password authenticated key exchange (aPAKE) that supports mutual authentication in a client-server setting without reliance on PKI and with security against pre-computation attacks upon server compromise. Prior aPAKE protocols did not use salt and if they did, the salt was transmitted in the clear from server to user allowing for the building of targeted pre-computed dictionaries. OPAQUE security has been proven by Jarecki et al. (Eurocrypt 2018) in a strong and universally composable formal model of aPAKE security. In addition, the protocol provides forward secrecy and the ability to hide the password from the server even during password registration.

Strong security, versatility through modularity, good performance, and an array of additional features make OPAQUE a natural candidate for practical use and for adoption as a standard. To this end, this draft presents several optimized instantiations of OPAQUE and ways of integrating OPAQUE with TLS.

This draft presents a high-level description of OPAQUE highlighting its components and modular design. A detailed unambiguous specification for standardization will be presented in future revisions of this document, or separately.

Status of This Memo

This Internet-Draft is submitted in full conformance with the provisions of BCP 78 and BCP 79.

Internet-Drafts are working documents of the Internet Engineering Task Force (IETF). Note that other groups may also distribute working documents as Internet-Drafts. The list of current Internet-Drafts is

Copyright Notice

1. Introduction

Password authentication is the prevalent form of authentication in the web and in most other applications. In the most common implementation, a user authenticates to a server by entering its user id and password where both values are transmitted to the server under the protection of TLS. This makes the password vulnerable to TLS failures, including many forms of PKI attacks, certificate mishandling, termination outside the security perimeter, visibility to middle boxes, and more. Moreover, even under normal operation, passwords are always visible in plaintext form at the server upon TLS decryption (in particular, storage of plaintext passwords is not an uncommon security incident, even among security-conscious companies).

Asymmetric (or augmented) Password Authenticated Key Exchange (aPAKE) protocols are designed to provide password authentication and mutually authenticated key exchange without relying on PKI (except during user/password registration) and without disclosing passwords to servers or other entities other than the client machine. A secure aPAKE should provide the best possible security for a password protocol, namely, it should only be open to inevitable attacks: online impersonation attempts with guessed user passwords and offline dictionary attacks upon the compromise of a server and leakage of its password file. In the latter case, the attacker learns a mapping of a user's password under a one-way function and uses such a mapping to validate potential guesses for the password. Crucially important is for the password protocol to use an unpredictable one-way mapping or otherwise the attacker can pre-compute a deterministic list of mapped passwords leading to almost instantaneous leakage of passwords upon server compromise.

Quite surprisingly, in spite of the existence of multiple designs for (PKI-free) aPAKE protocols, none of these protocols is secure against pre-computation attacks. In particular, none of these protocols can use the standard technique against pre-computation that combines *secret* random values (salt) into the one-way password mappings. Either these protocols do not use salt at all or, if they do, they transmit the salt from server to user in the clear, hence losing the secrecy of the salt and its defense against pre-computation. Furthermore, the transmission of salt may incur additional protocol messages.

This draft describes OPAQUE, a PKI-free secure aPAKE that is secure against pre-computation attacks and capable of using secret salt. OPAQUE has been recently defined and studied by Jarecki et al. [OPAQUE] who prove the security of the protocol in a strong aPAKE model that ensures security against pre-computation attacks and is formulated in the Universal Composability framework [Canetti01] under the random oracle model. In contrast, very few aPAKE protocols have been proven formally and those proven were analyzed in a weak security model that allows for pre-computation attacks (e.g., [GMR06]). This is not just a formal issue: these protocols are actually vulnerable to such attacks! Furthermore, as far as we know, protocols discussed recently as candidates for standardization (e.g., SPAKE2+ [I-D.irtf-cfrg-spake2] and AugPAKE [RFC6628]) do not enjoy a proof of security, not even in a weak model. The same holds for the SRP protocol [RFC2945]. VTBPEKE is analyzed in [VTBPEKE] in a weak model that allows for pre-computation attacks and, as in all the above cases (except OPAQUE), does not accommodate secret salt.

OPAQUE's design builds on a line of work initiated in the seminal paper of Ford and Kaliski [FK00] and is based on the HPAKE protocol of Xavier Boyen [Boyen09] and the (1,1)-PPSS protocol from Jarecki et al. [JKKX16]. None of these papers considered security against pre-computation attacks or presented a proof of aPAKE security (not even in a weak model).

In addition to its proven resistance to pre-computation attacks, OPAQUE's security features include forward secrecy (essential for protecting past communications in case of password leakage) and the ability to hide the password from the server - even during password registration. Moreover, good performance and an array of additional features make OPAQUE a natural candidate for practical use and for adoption as a standard. Such features include the ability to increase the difficulty of offline dictionary attacks via iterated hashing or other hardening schemes, and offloading these operations to the client (that also helps against online guessing attacks); extensibility of the protocol to support storage and retrieval of user's secrets solely based on a password; and being amenable to a multi-server distributed implementation where offline dictionary attacks are not possible without breaking into a threshold of servers (such distributed solution requires no change or awareness on the client side relative to a single-server implementation).

OPAQUE is defined and proven as the composition of two functionalities: An Oblivious PRF (OPRF) and a key-exchange protocol. It can be seen as a "compiler" for transforming any key-exchange protocol (with KCI security and forward secrecy - see below) into a secure aPAKE protocol. In OPAQUE, the user stores a secret private key at the server during password registration and retrieves this key each time it needs to authenticate to the server. The OPRF security properties ensure that only the correct password can unlock the private key while at the same time avoiding potential offline guessing attacks. This general composability property provides great flexibility and enables a variety of OPAQUE instantiations, from optimized performance to integration with TLS. The latter aspect is of prime importance as the use of OPAQUE with TLS constitutes a major security improvement relative to the standard password-over-TLS practice. At the same time, the combination with TLS builds OPAQUE as a fully functional secure communications protocol and can help provide privacy to account information sent by the user to the server prior to authentication.

The KCI property required from KE protocols for use with OPAQUE states that knowledge of a party's private key does not allow an attacker to impersonate others to that party. This is an important security property achieved by most public-key based KE protocols, including protocols that use signatures or public key encryption for authentication. It is also a property of many implicitly authenticated protocols (e.g., HMQV) but not all of them. We also note that key exchange protocols based on shared keys do not satisfy the KCI requirement, hence they are not considered in the OPAQUE setting. We note that KCI is needed to ensure a crucial property of OPAQUE: even upon compromise of the server, the attacker cannot impersonate the user to the server without first running an exhaustive dictionary attack. Another essential requirement from KE protocols for use in OPAQUE is to provide forward secrecy (against active attackers).

This draft presents a high-level description of OPAQUE highlighting its components and modular design. A detailed unambiguous specification for standardization will be presented in future revisions of this document, or separately.

We describe OPAQUE with a specific instantiation of the OPRF component over elliptic curves and with a few KE schemes, including the HMQV [HMQV] and SIGMA [SIGMA] protocols. We also present several strategies for integrating OPAQUE with TLS 1.3 [RFC8446] offering different tradeoffs between simplicity, performance and user privacy. See also the companion draft [I-D.sullivan-tls-opaque]. In general, the modularity of OPAQUE's design makes it easy to integrate with additional key-exchange protocols, e.g., IKEv2.

The computational cost of OPAQUE is determined by the cost of the OPRF, the cost of a regular Diffie-Hellman exchange, and the cost of authenticating such exchange. In our elliptic-curve implementation of the OPRF, the cost for the client is two exponentiations (one or two of which can be fixed base) and one hashing-into-curve operation [I-D.irtf-cfrg-hash-to-curve]; for the server, it is just one exponentiation. The cost of a Diffie-Hellman exchange is as usual two exponentiations per party (one of which is fixed-base). Finally, the cost of authentication per party depends on the specific KE protocol: it is just 1/6 of an exponentiation with HMQV and it is one signature generation and verification in the case of SIGMA and TLS 1.3. These instantiations preserve the number of messages in the underlying KE protocol except in one of the TLS instantiations where user privacy may require an additional round trip (this can be saved if using a mechanism similar to the proposed ESNI extension [I-D.ietf-tls-esni]).

1.1. Terminology

In this document, the key words "MUST", "MUST NOT", "REQUIRED", "SHALL", "SHALL NOT", "SHOULD", "SHOULD NOT", "RECOMMENDED", "MAY", and "OPTIONAL" are to be interpreted as described in BCP 14, RFC 2119 [RFC2119]

1.2. Notation

Throughout this document the first argument to a keyed function represents the key; separated by a

semicolon are the function inputs typically implemented as an unambiguous concatenation of strings (details of encodings are left for a future, more detailed specification).

Except if said otherwise, random choices in this specification refer to drawing with uniform distribution from a given set (i.e., "random" is short for "uniformly random").

The name OPAQUE: A homonym of O-PAKE where O is for Oblivious (the name OPAKE was taken).

2. DH-OPRF

A fundamental piece in the definition of OPAQUE is an Oblivious Pseudo Random Function (OPRF).

An Oblivious PRF (OPRF) is an interactive protocol between a server S and a user U defined by a special pseudorandom function (PRF), denoted F. The server's input to the protocol is a key k for PRF F and the user's input is a value x in the domain of F. At the end of the protocol, U learns $F(k;x)$ and nothing else while S learns nothing from the protocol execution (in particular nothing about x or the value $F(k;x)$).

OPAQUE uses a specific OPRF instantiation, called DH-OPRF, where the PRF, denoted F, is defined as follows (a detailed specification appears in [I-D.sullivan-cfrg-voprf] which defines several instantiation suites for DH-OPRF).

Parameters: Hash function H (e.g., SHA2 or SHA3 function) with 256-bit output at least, a cyclic group G of prime order q, a generator g of G, and hash function H' mapping arbitrary strings into G (where H' is modeled as a random oracle).

- DH-OPRF domain: Any string
- DH-OPRF range: The range of the hash function H
- DH-OPRF key: A random element k in [0..q-1]; denote $v=g^k$
- DH-OPRF Operation: $F(k; x) = H(x, v, H'(x)^k)$

Protocol for computing DH-OPRF, U with input x and S with input k:

- U: choose random r in [0..q-1], send alpha=$H'(x)*g^r$ to S
- S: upon receiving a value alpha, respond with $v=g^k$ and beta=alpha^k
- U: upon receiving values beta and v, set the PRF output to
  $H(x, v, beta*v^{-r})$

All received values (alpha, beta, v) are checked to be elements in G other than the identity. A party aborts if the check fails. In the case of Elliptic Curves this test is typically inexpensive - see [I-D.irtf-cfrg-spake2] for ways to deal with this check (including co-factor exponentiation) that apply to DH-OPRF as well.

Note (exponential blinding): An alternative way of computing DH-OPRF is for U to send alpha= $(H'(x))^r$ in the first message and set the function output to $H(x,v,beta^{1/r})$ upon receiving S's response. However, note that the multiplicative blinding above allows for a more efficient implementation as the g^r exponentiation uses a fixed base. Moreover, in cases where the user caches v (e.g., for sites it visits often) then one can also use a fixed-base optimized computation of the exponentiation v^{-r}. In any case, it is up to the user to decide on what form of blinding it uses; the server's response is the same in both cases.

Note: For elliptic curve implementations of DH-OPRF, the hashing into the curve operation has been studied extensively with known efficient implementations, see [I-D.irtf-cfrg-hash-to-curve] and references therein.

2.1. Hardening OPRF via user iterations

Protocol OPAQUE can be further strengthened against offline dictionary attacks by applying to the output of DH-OPRF an iterated hash for some number n of iterations. This increases the cost of an offline attack upon the compromise of the server as the attacker will need to perform n iterations for each guess of PwdU it tries to validate. For this purpose we would re-define DH-OPRF as $F(k;x) = I^n(H(x, v, H'(x)^k))$ where I is an appropriately chosen function and the symbol I^n denotes n iterations of function I. More generally, any form of hardened password-based key derivation can be used, e.g., PBKDF2 [RFC8018], Argon 2 [I-D.irtf-cfrg-argon2], scrypt [RFC7914]. As we will see, in OPAQUE it is the user who runs this key derivation. The iteration value n or any other parameter required by these functions can be set as public constants or can be set at time of password registration and later communicated by the server as part of its OPAQUE message. (Note: these hardened KDFs often require a salt value as input; for OPAQUE this value can be be set to a constant, e.g., all zeros.)

3. OPAQUE Specification

OPAQUE consists of the concurrent run of an OPRF protocol and a key-exchange protocol KE (one that provides mutual authentication based on public keys and satisfies the KCI requirement). We first define OPAQUE in a generic way based on any OPRF and any PK-based KE, and later show specific instantiation using DH-OPRF (defined in Section 2) and several KE protocols. The user takes the role of initiator in these protocols and the server the responder's. The private-public keys for the user are denoted PrivU and PubU, and for the server PrivS and PubS. In Section 3.4 we augment the protocol with user-side hardening to increase resistance to offline attacks in case of server compromise.

3.1. Password registration

Password registration is run between a user U and a server S. It is assumed that the user can authenticate the server during this registration phase (this is the only part in OPAQUE that requires some form of authenticated channel, either physical, out-of-band, PKI-based, etc.)

- U chooses password PwdU and a pair of private-public keys PrivU and PubU for the given protocol KE.
- S chooses OPRF key kU (random and independent for each user U) and sets $vU = g^{kU}$; it also chooses its own pair of private-public keys PrivS and PubS for use with protocol KE (the server can use the same pair of keys with multiple users), and sends PubS to U.
- U and S run OPRF(kU;PwdU) as defined in Section 2 with only U learning the result, denoted RwdU (mnemonics for "Randomized PwdU").
- U generates an "envelope" EnvU defined as

  EnvU = AuthEnc(RwdU; PrivU, PubU, PubS)

  where AuthEnc is an authenticated encryption function with the "random-key robustness" property as specified in Section 3.1.1 below. In EnvU, all values require authentication and PrivU also requires encryption. However, for simplicity and to hide the EnvU contents, we specify that all values are encrypted (not just authenticated). PubU can be omitted from EnvU if it is not needed for running the key-exchange protocol by the client or if it can be reconstructed from PrivU.
- U sends EnvU and PubU to S and erases PwdU, RwdU and all keys.
  S stores (EnvU, PubS, PrivS, PubU, kU, vU) in a user-specific record. If PrivS and PubS are used for multiple users, S can store these values separately and omit them from the user's record.

Note (salt). We note that in OPAQUE the OPRF key acts as the secret salt value that ensures the infeasibility of pre-computation attacks. No extra salt value is needed.

Note (password rules). The above procedure has the significant advantage that the user's password is not disclosed to the server even during registration. Some sites require learning the user's password for enforcing password rules. Doing so voids this important security property of OPAQUE and is not recommended. Moving the password check procedure to the client side is a more secure alternative (limited checks at the server are possible to implement, e.g., detecting repeated passwords).

3.1.1. Implementing the EnvU envelop

The function AuthEnc used to compute EnvU needs to satisfy a property called "random-key robustness". That is, given a pair of random AuthEnc keys, it should be infeasible to create an authenticated ciphertext that successfully decrypts under the two keys. Not all AuthEnc schemes enjoy this property, and this includes the standard GCM mode. We describe a few methods for implementing a key-robust AuthEnc scheme (a default mechanism needs to be chosen when specifying a standard).

(1) Encrypt-then-HMAC: Implement the AuthEnc scheme using any encryption function in encrypt-then-pad-then-MAC mode where the MAC is implemented with HMAC with a tag size of at least 256 bits (HMAC ensures robustness through the collision-resistant property of the underlying hash function). This requires two separate keys, one for encryption and one for HMAC, which can be derived from RwdU using, for example, the HKDF-Expand function from [RFC5869].

[Author note: Is there a standardized definition of Encrypt-then-HMAC that one can refer? Two possible sources are RFC 7366 and signal [SIGNAL]. Note that here we just need it for generating EnvU, without the complexities of a full interactive secure-channel implementation.]

(2) Fixed-string MAC: Any AuthEnc scheme can be made random-key robust by simply concatenating to it a MAC computed on a fixed string. Specifically, one derives two keys from RwdU: one is a key to the AuthEnc function and the other is a key to a MAC (any secure MAC function works). EnvU is then

defined as before except that to the output of AuthEnc one concatenates the output of the MAC function computed with the second derived key on a fixed string (e.g., all zeros).

(3) To make GCM key robust one can use any of the two methods above or use the following GCM-specific mechanism. To the plaintext specified by the definition of EnvU, concatenate a string of all zeros so that the last full block of the resultant plaintext is an all-zero block (of the length of the underlying block cipher, i.e., 128 for AES). Use regular GCM on this augmented plaintext with a nonce value of zero (a zero nonce is OK since RwdU is a one-time encryption key, namely, no two EnvU encryptions will use the same RwdU, except for negligible probability).

Note: With the 128-bit output for the authentication tag in GCM-AES, there is a 2^{64} collision attack against the robustness property that seems impractical in the OPAQUE setting but should be considered. The first two methods above don't have this issue.

Note (authentication-only EnvU): It is possible to dispense with encryption in the construction of EnvU to obtain a shorter EnvU (resulting in less storage at the server and less communication from server to client). In this case, p_u is derived using a PRF keyed with RwdU. For cases where p_u is not a random string of a given length, e.g., RSA, we define a more general procedure. Namely, what's derived from RwdU is a random seed used as an input to a key generation procedure that generates (p_u, P_U). The resultant scheme is as follows:

- Set Km = f(0), seed = f(1), (p_u, P_u) = KeyGen(seed), where f is a regular PRF keyed with RwdU (it can also be implemented with HKDF-Expand).
- Set EnvU = (P_s, HMAC_Km(P_s)) where P_s is the server's public key. HMAC_Km denotes HMAC computed with key Km. This MAC needs to have a key-robust property hence it is implemented with HMAC that satisfies this property.

Assuming that the server derives per-user kU using a PRF and global key, and that it uses the same pair (p_s,P_s) with all users, the per-user storage for the server consists of P_u and HMAC_Km(P_s), a total of 64-byte overhead with 256-bit curve and hash. Yet, in spite of these storage and communication savings, using encryption for EnvU avoids the need for running the key generation procedure at the client and allows OPAQUE to also serve as a retrieval mechanism for other secrets or credentials.

3.2. Online OPAQUE protocol

After registration, the user and server can run the OPAQUE protocol as a password-authenticated key exchange. The protocol proceeds as follows:

- User transmits user/account information to the server so that the server can retrieve the user's record.
- Server sends EnvU and vU to user.
- Server and User execute the OPRF protocol as defined in Section 2; User sets RwdU to the result of this computation.
- User decrypts EnvU using RwdU to obtain PrivU, PubU, PubS.
- User and server run the specified KE protocol using their respective public and private keys.

Note that the steps preceding the run of KE can be arranged in just two messages (one from the user and a response from the server). Furthermore, OPAQUE is optimized by running the OPRF and KE concurrently with interleaved and combined messages (while preserving the internal ordering of messages in each protocol). In all cases, the user needs to obtain EnvU, vU and RwdU (i.e., complete the OPRF protocol) before it can use its own private key PrivU and the server's public key PubS in the run of KE.

3.3. OPAQUE Instantiations

We present several instantiations of OPAQUE using DH-OPRF as the OPRF and different KE protocols. For the sake of concreteness we focus on KE protocols consisting of three messages, denoted KE1, KE2, KE3, and such that KE1 and KE2 include DH values sent by user and server, respectively, and KE3 provides explicit user authentication. (As shown in [OPAQUE], OPAQUE cannot use less than three messages so the 3-message instantiations presented here are optimal in terms of number of messages.)

Generic OPAQUE with 3-message KE:

- C to S: Uid, alpha=H'(PwdU)*g^r, KE1
- S to C: beta=alpha^kU, vU, EnvU, KE2
- C to S: KE3

Key derivation and other details of the protocol are fully specified by the KE scheme. We do note that by the results in [OPAQUE], KE2 and KE3 should include authentication of the OPRF messages (or at least of the value alpha).

We provide two instantiations of OPAQUE (with HMQV and SIGMA-I) next and discuss integration with TLS in [RFC8446].

3.3.1. Instantiation with HMQV

The integration of OPAQUE with HMQV [HMQV] leads to the most efficient instantiation of OPAQUE in terms of exponentiation count. Performance is close to optimal due to the low cost of authentication in HMQV: Just 1/6 of an exponentiation for each party over the cost of a regular DH exchange. The private and public keys of the parties are Diffie-Hellman keys, namely, PubU=g^PrivU and PubS=g^PrivS. The HMQV exchange can be represented schematically as follows:

- KE1 = g^x
- KE2 = g^y, Mac(Km1; g^x, g^y)
- KE3 = Mac(Km2; g^y, g^x)

Keys in HMQV, namely, MAC keys Km1, Km2 and session/traffic keys are derived from a common key K computed as follows:

C computes K = H((g^y * PubS^e)^{x + d*PrivU))

S computes K = H((g^x * PubU^d)^{y + e*PrivS))

where d = H(g^x, IdS) and e = H(g^y, IdU), and IdU, IdS represent the identities of user and server (typically, a user id and domain name, respectively, and can include the party's public key value). The function H can be the same as used for the DH-OPRF computation. Using multi-exponentiation optimization, the computation of K involves a single multi-exponentiation whose cost is only 17% more than a regular exponentiation.

This is a minimal skeleton. A fully-specified protocol will include additional details and a careful key derivation scheme. In particular, the Mac computation will cover the whole preceding transcript. In addition, the parties will check group membership for g^x, g^y or use co-factor computation, e.g., as in [I-D.irtf-cfrg-spake2], and also reject these values if they equal the identity. The check for PubU and PubS can be done only once at user registration.

Note (IP disclosure): IBM has a patent that covers HMQV. If there is interest in standardizing this mode one can check if a free license of the HMQV patent could be provided for such use.

3.3.2. Instantiation with SIGMA-I

We show how OPAQUE can be built around the 3-message SIGMA-I protocol [SIGMA]. This example is significant as it shows integration with a signature-based KE protocol and because TLS 1.3 follows the design of SIGMA-I hence the example helps understanding the proposed integration of OPAQUE with TLS in [RFC8446].

SIGMA-I can be represented schematically as follows:

- KE1 = g^x
- KE2 = g^y, Sig(PrivS; g^x, g^y), Mac(Km1; IdS)
- KE3 = Sig(PrivU; g^y, g^x), Mac(Km2; IdU)

In this case, the private keys of user and server are signature keys. Key derivation is based on the DH value g^xy.

As before, this is only a skeleton to illustrate the protocol.
Full details need to be filled in for a self-contained specification. See also Section 4 for the integration of SIGMA in TLS 1.3.

3.4. Hardening OPAQUE via hardened key derivation

As noted in Section 2.1 one can increase the resistance of OPAQUE to offline attacks in case of server compromise by applying a password-hardening key derivation function (KDF) on top of the regular DH-OPRF when computing RwdU. Specifically, the user computes the OPRF in interaction with the server and inputs its result into the KDF in lieu of the password on which this function acts. The output of this computation is set as the value RwdU. Parameters to the KDF function can be set to public values or set at the time of password registration and stored at the server. In this case, the server communicates

these parameters to the user during OPAQUE executions together with the second OPRF message. We note that the salt value typically input into the KDF can be set to a constant, e.g., all zeros.

4. Integrating OPAQUE with TLS 1.3

Note: This section is intended as a basis for discussion on ways to integrate OPAQUE with TLS (particularly TLS 1.3). Precise protocol details are left for a future specification. A preliminary draft is [I-D.sullivan-tls-opaque].

As stated in the introduction, the standard password-over-TLS mechanism for password authentication suffers from significant weaknesses due to the essential reliance of the protocol on PKI and the exposure of passwords to the server (and other observers) upon TLS decryption. Here we propose integrating OPAQUE with TLS in order to remove these vulnerabilities while at the same time armoring TLS itself against PKI failures. Such integration also benefits OPAQUE by leveraging the standardized negotiation and record-layer security of TLS. Furthermore, TLS can offer an initial PKI-authenticated channel to protect the privacy of account information such as user name transmitted between client and server.

If one is willing to forgo protection of user account information transmitted between user and server, integrating OPAQUE with TLS 1.3 is relatively straightforward and follows essentially the same approach as with SIGMA-I in Section 3.3.2. Specifically, one reuses the Diffie-Hellman exchange from TLS and uses the user's private key PrivU retrieved from the server as a signature key for TLS client authentication. The integrated protocol will have as its first message the TLS's Client Hello augmented with user account information and with the DH-OPRF first message (the value alpha). The server's response includes the regular TLS 1.3 second flight augmented with the second OPRF message which includes the values beta, vU and EnvU. For its TLS signature, the server uses the private key PrivS whose corresponding public key PubS is authenticated as part of the user envelope EnvU (there is no need to send a regular TLS certificate in this case). Finally, the third flight consists of the standard client Finish message with client authentication where the client's signature is produced with the user's private key PrivU retrieved from EnvU and verified by the server using public key PubU.

The above scheme is depicted in Figure 1 where the sign + indicates fields added by OPAQUE, and DH-OPRF1, DH-OPRF2 denote the two DH-OPRF messages. Other messages in the figure are the same as in TLS 1.3. Notation {...} indicates encryption under handshake keys. Note that ServerSignature and ClientSignature are performed with the private keys defined by OPAQUE and they replace signatures by traditional TLS certificates.

```
    Client                                        Server

    ClientHello
    key_share
    + userid + DH−OPRF1      --------->
                                              ServerHello
                                               key_share
                                     {+ DH−OPRF2 + vU + EnvU}
                                          {+ ServerSignature}
                             <--------       {ServerFinished}

    {+ ClientSignature}
    {ClientFinished}         --------->
```

Figure 1: Integration of OPAQUE in TLS 1.3 (no userid confidentiality)

Note that in order to send DH-OPRF1 in the first message, the client needs to know the DH group the server uses for OPRF, or it needs to "guess" it. This issue already appears in TLS 1.3 where the client needs to guess the key_share group and it should be handled similarly in OPAQUE (e.g., the client may try one or more groups in its first message).

Protection of user's account information can be added through TLS 1.3 pre-shared/resumption mechanisms where the account information appended to the ClientHello message would be encrypted under the pre-shared key.

When a resumable session or pre-shared key between the client and the server do not exist, user account protection requires a server certificate. One option that does not add round trips is to use a mechanism similar to the proposed ESNI extension [I-D.ietf-tls-esni]. Without such extension, one would run a TLS 1.3 handshake augmented with two first OPAQUE messages interleaved between the second and third flight of the regular TLS handshake. That is, the protocol consists of five flights as follows: (i) A regular 2-flight 1-RTT handshake to produce handshake traffic keys authenticated by the server's TLS certificate; (ii) two messages that include user identification information, the DH-OPRF

messages exchanged between client and server, and the retrieved vU and EnvU, all encrypted under the handshake traffic keys (thus providing privacy to user account information); (iii) the TLS 1.3 client authentication flight where client authentication uses the user's private signature key PrivU retrieved from the server in step (ii).

Note that server authentication in (i) uses TLS certificates hence user account privacy (but not user authentication) is dependent on PKI. In cases where PKI authentication for the server is deemed acceptable then there is no need for further server authentication. However, if one wants to enforce server authentication without reliance on PKI, then the server needs to authenticate using the private key PrivS whose corresponding public key PubS is sent to the user as part of EnvU. There are two options: If PubS is the same as the public key the server used in the 1-RTT authentication (step (i)) then there is no need for further authentication. In this case, U gets assurance from the authenticated EnvU, not (only) from the PKI certificates. Otherwise, the server needs to send a signature under PrivS that is piggybacked to the second OPAQUE message in (ii). In this case the signature would cover the running transcript hash as is standard in TLS 1.3. The client signature in the last message also covers the transcript hash including the regular handshake and OPAQUE messages.

The above scheme is depicted in Figure 2. Please refer to the text before Figure 1 describing notation. Note the asterisk in the ServerSignature message. This indicates that this message is optional as it is used only if the server's key PubS in OPAQUE is different than the one in the server's certificate (transmitted in the second protocol flight).
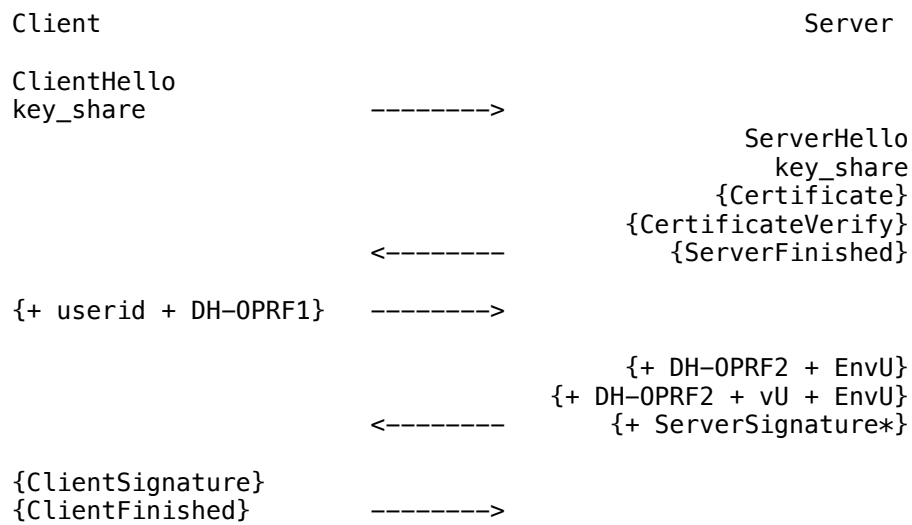
```
      Client                                              Server

      ClientHello
      key_share                 -------->
                                                       ServerHello
                                                        key_share
                                                      {Certificate}
                                                  {CertificateVerify}
                                <--------          {ServerFinished}

      {+ userid + DH-OPRF1}      -------->

                                                  {+ DH-OPRF2 + EnvU}
                                              {+ DH-OPRF2 + vU + EnvU}
                                <--------          {+ ServerSignature*}

      {ClientSignature}
      {ClientFinished}          -------->
```

Figure 2: Integration of OPAQUE in TLS 1.3 (with userid confidentiality)

We note that the above approaches for integration of OPAQUE with TLS may benefit from the post-handshake client authentication mechanism of TLS 1.3 and the exported authenticators from [I-D.ietf-tls-exported-authenticator]. Also, formatting of messages and negotiation information suggested in [I-D.barnes-tls-pake] can be used in the OPAQUE setting.

5. User enumeration

User enumeration refers to attacks where the attacker tries to learn whether a given user identity is registered with a server. Preventing such attack requires the server to act with unknown user identities in a way that is indistinguishable from its behavior with existing users. Here we suggest a way to implement such defense, namely, a way for simulating the values beta, vU and EnvU for non-existing users. Note that if the same pair of user identity UId and value alpha is received twice by the server, the response needs to be the same in both cases (since this would be the case for real users). For this, the server S will have two keys MK, MK' for a PRF f (this refers to a regular PRF such as HMAC or CMAC). Upon receiving a pair of user identity UId and value alpha for a non-existing user UId, S computes kU=f(MK; UId) and kU'=f(MK'; UId) and responds with three values beta=alpha^kU, vU=g^kU and EnvU where the latter is computed as the AuthEnc function with key kU' applied to the all-zero string (of the length of a regular EnvU plaintext). Care needs to be taken to avoid side channel leakage (e.g., timing) from helping differentiate these operations from a regular server response. The above requires changes to the server-side implementation but not to the protocol itself or the client side.

There is one form of leakage that the above allows and whose prevention would require a change in OPAQUE. Note that an attacker that tests a UId (and same alpha) twice and receives different responses can conclude that either the user registered with the service between these two activations or that the user was registered before but changed its password in between the activations (assuming

the server changes kU at the time of a password change). In any case, this indicates that UId is a registered user at the time of the second activation. To conceal this information, S can implement the derivation of kU as kU=f(MK; UId) also for registered users. Hiding changes in EnvU, however, requires a change in the protocol. Instead of sending EnvU as is, S would send an encryption of EnvU under a key that the user derives from the OPRF result (similarly to RwdU) and that S stores during password registration. During login, the user will derive this key from the OPRF result, use it to decrypt EnvU, and continue with the regular protocol. If S uses a randomized encryption, the encrypted EnvU will look each time as a fresh random string, hence S can simulate the encrypted EnvU also for non-existing users.

[Author note: How significant is the user enumeration issue? The first case above does not change the protocol so its implementation is a server's decision. The second case, however, requires a change in OPAQUE so it would be important to have feedback on whether this second order attack justifies such change.]

6. Security considerations

This is an early draft presenting the OPAQUE concept and its potential instantiations. More precise details and security considerations will be provided in future drafts. We note that the security of OPAQUE is formally proved in [OPAQUE] under a strong model of aPAKE security assuming the security of the OPRF function and of the underlying key-exchange protocol. In turn, the security of DH-OPRF is proven in the random oracle model under the One-More Diffie-Hellman assumption [JKKX16].

Best practices regarding implementation of cryptographic schemes apply to OPAQUE. Particular care needs to be given to the implementation of the OPRF regarding testing group membership and avoiding timing and other side channel leakage in the hash-to-curve mapping. Drafts [I-D.irtf-cfrg-hash-to-curve] and [I-D.sullivan-cfrg-voprf] have detailed instantiation and implementation guidance.

While one can expect the practical security of the OPRF function (namely, the hardness of computing the function without knowing the key) to be in the order of computing discrete logarithms or solving Diffie-Hellman, Brown and Gallant [BG04] and Cheon [Cheon06] show an attack that slightly improves on generic attacks. For the case that $q-1$ or $q+1$, where q is the order of the group G, has a t-bit divisor, they show an attack that calls the OPRF on 2^t chosen inputs and reduces security by $t/2$ bits, i.e., it can find the OPRF key in time $2^{q/2-t/2}$ and $2^{q/2-t/2}$ memory. For typical curves, the attack requires an infeasible number of calls and/or results in insignificant security loss [*]. Moreover, in the OPAQUE application, these attacks are completely impractical as the number of calls to the function translates to an equal number of failed authentication attempts by a *single* user. For example, one would need a billion impersonation attempts to reduce security by 15 bits and a trillion to reduce it by 20 bits - and most curves will not even allow for such attacks in the first place (note that this theoretical loss of security is with respect to computing discrete logarithms, not in reducing the password strength).

[*] Some examples (courtesy of Dan Brown): For P-384, 2^{90} calls reduce security from 192 to 147 bits; for NIST P-256 the options are 6-bit reduction with 2153 OPRF calls, about 14 bit reduction with 187 million calls and 20 bits with a trillion calls. For Curve25519, attacks are completely infeasible (require over 2^{100} calls) but its twist form allows an attack with 25759 calls that reduces security by 7 bits and one with 117223 calls that reduces security by 8.4 bits.

Note on user authentication vs. authenticated key exchange. OPAQUE provides PAKE (password-based authenticated key exchange) functionality in the client-server setting. While in the case of user identification, focus is often on the authentication part, we stress that the key exchange element is not less crucial. Indeed, in most cases user authentication is performed to enforce some policy, and the key exchange part is essential for binding this enforcement to the authentication step. Skipping the key exchange part is analogous to carefully checking a visitor's credential at the door and then leaving the door open for others to enter freely.

This draft complies with the requirements for PAKE protocols set forth in [RFC8125].

7. Acknowledgments

The OPAQUE protocol and its analysis is joint work of the author with Stas Jarecki and Jiayu Xu. This draft has benefited from comments by multiple people. Special thanks to Richard Barnes, Dan Brown, Eric Crockett, Fredrik Kuivinen, Kevin Lewi, Payman Mohassel, Jason Resch, Nick Sullivan.

8. References

8.1. Normative References

[RFC2119]   Bradner, S., "Key words for use in RFCs to Indicate Requirement Levels", BCP 14, RFC 2119, DOI 10.17487/RFC2119, March 1997.

8.2. Informative References

[BG04]   Brown, D. and R. Galant, "The static Diffie-Hellman problem", http://eprint.iacr.org/2004/306 , 2004.

[Boyen09]   Boyen, X., "HPAKE: Password authentication secure against cross-site user impersonation", Cryptology and Network Security (CANS) , 2009.

[Canetti01] Canetti, R., "Universally composable security: A new paradigm for cryptographic protocols", IEEE Symposium on Foundations of Computer Science (FOCS) , 2001.

[Cheon06]   Cheon, J., "Security analysis of the strong Diffie-Hellman problem", Euroctypt 2006 , 2006.

[FK00]   Ford, W. and B. Kaliski, Jr, "Server-assisted generation of a strong secret from a password", WETICE , 2000.

[GMR06]   Gentry, C., MacKenzie, P. and . Z, Ramzan, "A method for making password-based key exchange resilient to server compromise", CRYPTO , 2006.

[HMQV]   Krawczyk, H., "HMQV: A high-performance secure Diffie-Hellman protocol", CRYPTO , 2005.

[I-D.barnes-tls-pake]   Barnes, R. and O. Friel, "Usage of PAKE with TLS 1.3", Internet-Draft draft-barnes-tls-pake-04, July 2018.

[I-D.ietf-tls-esni]   Rescorla, E., Oku, K., Sullivan, N. and C. Wood, "Encrypted Server Name Indication for TLS 1.3", Internet-Draft draft-ietf-tls-esni-04, July 2019.

[I-D.ietf-tls-exported-authenticator]   Sullivan, N., "Exported Authenticators in TLS", Internet-Draft draft-ietf-tls-exported-authenticator-09, May 2019.

[I-D.irtf-cfrg-argon2]   Biryukov, A., Dinu, D., Khovratovich, D. and S. Josefsson, "The memory-hard Argon2 password hash and proof-of-work function", Internet-Draft draft-irtf-cfrg-argon2-08, October 2019.

[I-D.irtf-cfrg-hash-to-curve]   Faz-Hernandez, A., Scott, S., Sullivan, N., Wahby, R. and C. Wood, "Hashing to Elliptic Curves", Internet-Draft draft-irtf-cfrg-hash-to-curve-04, July 2019.

[I-D.irtf-cfrg-spake2]   Ladd, W. and B. Kaduk, "SPAKE2, a PAKE", Internet-Draft draft-irtf-cfrg-spake2-09, October 2019.

[I-D.sullivan-cfrg-voprf]   Davidson, A., Sullivan, N. and C. Wood, "Oblivious Pseudorandom Functions (OPRFs) using Prime-Order Groups", Internet-Draft draft-sullivan-cfrg-voprf-03, March 2019.

[I-D.sullivan-tls-opaque]   Sullivan, N., Krawczyk, H., Friel, O. and R. Barnes, "Usage of OPAQUE with TLS 1.3", Internet-Draft draft-sullivan-tls-opaque-00, March 2019.

[JKKX16]   Jarecki, S., Kiayias, A., Krawczyk, H. and J. Xu, "Highly-efficient and composable password-protected secret sharing (or: how to protect your bitcoin wallet online)", IEEE European Symposium on Security and Privacy , 2016.

[OPAQUE]   Jarecki, S., Krawczyk, H. and J. Xu, "OPAQUE: An Asymmetric PAKE Protocol Secure Against Pre-Computation Attacks", Eurocrypt , 2018.

[RFC2945]   Wu, T., "The SRP Authentication and Key Exchange System", RFC 2945, DOI 10.17487/RFC2945, September 2000.

[RFC5869]   Krawczyk, H. and P. Eronen, "HMAC-based Extract-and-Expand Key Derivation Function (HKDF)", RFC 5869, DOI 10.17487/RFC5869, May 2010.

[RFC6628]   Shin, S. and K. Kobara, "Efficient Augmented Password-Only Authentication and Key Exchange for IKEv2", RFC 6628, DOI 10.17487/RFC6628, June 2012.

[RFC7914]   Percival, C. and S. Josefsson, "The scrypt Password-Based Key Derivation Function", RFC 7914, DOI 10.17487/RFC7914, August 2016.

[RFC8018]   Moriarty, K., Kaliski, B. and A. Rusch, "PKCS #5: Password-Based Cryptography Specification Version 2.1", RFC 8018, DOI

10.17487/RFC8018, January 2017.

[RFC8125]   Schmidt, J., "Requirements for Password-Authenticated Key Agreement (PAKE) Schemes", RFC 8125, DOI 10.17487/RFC8125, April 2017.

[RFC8446]   Rescorla, E., "The Transport Layer Security (TLS) Protocol Version 1.3", RFC 8446, DOI 10.17487/RFC8446, August 2018.

[SIGMA]   Krawczyk, H., "SIGMA: The SIGn-and-MAc approach to authenticated Diffie-Hellman and its use in the IKE protocols", CRYPTO , 2003.

[SIGNAL]   "Signal recommended cryptographic algorithms", https://signal.org/docs/specifications/doubleratchet/#recommended-cryptographic-algorithms , 2016.

[VTBPEKE]   Pointcheval, D. and G. Wang, "Verifier-based Two-Basis Password Exponential Key Exchange", AsiaCCS , 2017.